

How Sentinel Q gets its news

Step 2 (Ingest) methodology · free sources only · nothing billed · no API key · articles read in memory, never saved

1. Principle

For each stock we take the **latest 100 articles whose headline names the company**, newest first, looking back at most 12 months. We keep scraping further back until we have 100 or the 12-month limit is reached. The model never searches, browses or recalls anything; it only labels what we hand it. Every article keeps its URL and date, and anything dropped is listed with a reason.

2. Sources (free)

Source	Used for	Notes
Google News RSS (default, India edition)	Live news, latest-first	Query per company with after :/before : date operators. Index is live, so it is not a point-in-time archive.
GDELT DOC 2.0	Dated archive for past as-of dates	Free; each request only reaches ~3 months, so windows are rolled back. Chosen automatically (--news auto) when the as-of date is in the past.
Verified events CSV	Human-checked filings	Events a person confirmed against a filing are merged before scoring and flagged as verified.

3. The rolling-window search (per stock)

- **Query:** the company's name plus up to 3 aliases from `portfolio/universe.csv`, as an OR-query in quotes, e.g. ("Titan Company" OR "Titan") after:2026-06-01 before:2026-07-04.
- **Windows roll backwards, newest first:** last 30 days, then days 31-90, 91-180, 181-365. Each is a new, non-overlapping query; we stop as soon as 100 usable titles are in hand.
- **Usable = the headline names the company.** Loose hits (e.g. a ceiling-fan story for a similarly named firm) do not count towards the 100. Syndicated copies of the same headline are counted once.
- **Selection:** the latest 100 usable headlines, newest first. Other articles only fill the list if fewer than 100 usable ones exist.
- **Optional body text:** with --fetch-text the page is read in memory (25 s budget, first paragraphs only) for better labels; otherwise the model labels headline + snippet.

4. Three passes, so that nothing important hides behind the latest 100

- **Sentiment pass** - the latest-100 search above. This is the only pass that feeds company sentiment.
- **Results (fundamentals) pass** - a separate 12-month search for results, profit, revenue, guidance, order wins. It anchors sentiment so a quiet recent quarter cannot hide a year of results.
- **Governance pass** - a separate 12-month search for SEBI/RBI/penalty/probe/resignation/auditor terms (up to 30 per stock), because governance events are rare and often older than the latest 100.

5. Reliability and honesty rules

- **Rate limits:** flat 20 s waits and up to 3 retries; a health check before the run; the run stops after 3 consecutive stocks fail rather than silently scoring empty data.
- **Minimum evidence:** fewer than 8 relevant articles, or no results print, means sentiment is shown as **n/a (Insufficient Data)**, never a guessed number.
- **Time travel guard:** a live index with a past as-of date is refused unless explicitly allowed, and then stamped "not a point-in-time backtest".
- **Nothing stored:** articles stream through memory (scrape stock N while stock N-1 is labelled). Only the audit record is written: URL, date, headline, label, drop reason.
- **Prices never enter a score.** Share-price headlines are typed price_move and excluded from sentiment.

6. How to check it yourself

`python -m sentinelq inspect news --stock TITAN` shows what the search returns for one stock. After a run, `audit/queries.jsonl` lists every query and window with the count returned, `dropped.jsonl` every article dropped and why, and `work/run.log` the live trace. Code: `sentinelq/ingest/gnews.py`, `gdelit.py`, `select.py`.